

hw5

October 5, 2023

1 Homework 5

1.1 Question 1 (5 points)

- Read the `mpg` dataset in a `pandas` data frame.
- Execute the `describe` method of the data frame you created.
- What function in R corresponds to `describe`?

```
[3]: #!/pandoc --version
```

```
[77]: import pandas as pd
mpg=pd.read_csv('mpg.csv')
mpg.describe()
```

```
[77]:
```

	displ	year	cyl	cty	hwy
count	234.000000	234.000000	234.000000	234.000000	234.000000
mean	3.471795	2003.500000	5.888889	16.858974	23.440171
std	1.291959	4.509646	1.611534	4.255946	5.954643
min	1.600000	1999.000000	4.000000	9.000000	12.000000
25%	2.400000	1999.000000	4.000000	14.000000	18.000000
50%	3.300000	2003.500000	6.000000	17.000000	24.000000
75%	4.600000	2008.000000	8.000000	19.000000	27.000000
max	7.000000	2008.000000	8.000000	35.000000	44.000000

1.1.1 What function in R corresponds to `describe`?

Summary built in function

1.2 Question 2 (25 points)

This will make you miss the `ggplot2` library! :D

Modify the following plot

```
_ = sns.lmplot(
    data = mpg,
    x = "displ",
    y = "hwy",
    ci = None,
    line_kws = {"color": "black"},
```

```

        scatter_kws = {"s": 10},
        lowess = True,
        x_jitter = 0.5,
        y_jitter = 0.5
    ).set(
        title = "Fuel consumption by engine displacement",
        xlabel = "Engine displacement",
        ylabel = "Fuel consumption"
    ).tight_layout()

```

in such a way that:

1. the points in the plot are colored by the variable `class` in the `mpg` dataset
2. there is a single trend line applied that applies to all points in the plot instead of one trend line for each group defined by the `class` variable in the `mpg` dataset.

Hint: - First, create a `lmplot` and set in that plot `hue = "class"` (to color the points by the “class” variable) and `fit_reg = False` (to suppress the trend line); assign the output of the `lmplot` function to a variable named `lm_plot` (you’ll need this in a second). - Next, in the same cell, create a `regplot` and set `scatter = False` (to suppress plotting the points, which are already produced by the previous plot) as well as `ax = lm_plot.ax` (to tell `seaborn` to overlay the new plot to the previous one, rather than creating a new plot). - You should still pass relevant parameters to each of the two components of the plot (`lmplot` and `regplot`) in a sensible way so to get the desired final plot as the result.

Here’s a “skeleton” of the code that you’ll need. It’s up to you to fill it out!

```

lm_plot = sns.lmplot(
    ... <- configure the lmplot appropriately
).set(
    title = "Fuel consumption by engine displacement",
    xlabel = "Engine displacement",
    ylabel = "Fuel consumption"
).tight_layout()

_ = sns.regplot(
    ... <- configure the regplot appropriately
)

```

```


[78]: import seaborn as sns
import pandas as pd
import matplotlib.pyplot as plt
lm_plot = sns.lmplot(
    data = mpg,
    x = "displ",
    y = "hwy",
    ci = None,
    line_kws = {"color": "black"},
    scatter_kws = {"s": 10},

```

```

    hue="class",
    fit_reg=False,
    lowess = True,
    x_jitter = 0.5,
    y_jitter = 0.5
).set(
    title = "Fuel consumption by engine displacement",
    xlabel = "Engine displacement",
    ylabel = "Fuel consumption"
).tight_layout()
s=sns.regplot(
    data = mpg,
    x = "displ",
    y = "hwy",
    scatter=False,
    line_kws = {"color": "black"},
    ax=lm_plot.ax,ci=None
)
plt.show()

```



```
[ ]: #!pip install seaborn==0.12.2
```

1.3 Question 3 (20 points)

Use the plotnine library to recreate this ggplot2 plot (from lab 1) in Python:

```
ggplot(data = diamonds, mapping = aes(x = color, fill = clarity)) +  
  geom_bar(position = "dodge") +  
  scale_fill_brewer(palette = "Set2") +  
  labs(  
    x = "Color",  
    y = element_blank(),  
    fill = "Diamond clarity",  
    title = "Clarity distribution by color"  
  )
```


Hint: [this stackoverflow](#) discussion might come in handy to set the palette colors. Also, note that the variable `clarity` is actually an ordered categorical variable (to see this, read the docs in R for the `diamonds` dataset with `? diamonds`). We can create a `pandas` equivalent of an R factor by using the `pandas.Categorical` class (see the docs [here](#)). To get the same ordering that we got in R for the dodged bars, you will need to transform the `clarity` variable in the `diamonds` dataset that you read with `pandas` to a `pandas.Categorical` variable.

```
[4]: #!pip install plotnine
```

```
[6]: from plotnine import *  
import pandas as pd
```

```
[7]: diamonds=pd.read_csv('diamonds.csv')
```

```
[24]: # Define the desired order for clarity levels  
clarity_order = ['I1', 'SI2', 'SI1', 'VS2', 'VS1', 'VVS2', 'VVS1', 'IF']  
diamonds['clarity'] = pd.Categorical(diamonds['clarity'],  
    ↪categories=clarity_order, ordered=True)  
  
(ggplot(data = diamonds, mapping = aes(x = 'color', fill = 'clarity')) +  
  geom_bar(position = "dodge") +  
  scale_fill_brewer(type="qual", palette = "Set2") +  
  labs(  
    x = "Color",  
    y = element_blank(),  
    fill = "Diamond clarity",  
    title = "Clarity distribution by color"  
  ))
```



[24]: <Figure Size: (640 x 480)>

```
[6]: #!pip install plotnine
```

1.4 Question 4 (25 points)

Solve again exercises 1 and 4 of Homework 2, but this time use `pandas`!

You are encouraged to use the published solutions of Homework 2 and “translate” them from R’s `tidyverse`/`dplyr` to Python’s `pandas`.

You might find this [“translation guide”](#) helpful!

```
[40]: # Load the flights and airlines datasets (assuming you have them as CSV files)
flights = pd.read_csv('flights.csv') # Replace with your dataset filename
airlines = pd.read_csv('airlines.csv') # Replace with your dataset filename

# Filter flights that occurred between July and November
flights = flights[(flights['month'] >= 7) & (flights['month'] <= 11)]

# Calculate the gain
flights['gain'] = flights['arr_delay'] - flights['dep_delay']
```

```

# Group by airline and compute average gain (ignoring NAs)
#Note: missing values are automatically skipped when calculating the mean. You
↳do not need to specify a parameter to skip them.
average_gain = flights.groupby('carrier')['gain'].mean().reset_index()

# Sort by average_gain in decreasing order
average_gain = average_gain.sort_values(by='gain', ascending=False)

# Join with the airlines dataset
average_gain = average_gain.merge(airlines, left_on='carrier',
↳right_on='carrier', how='left')

result_tibble = average_gain[['gain', 'name']]
# Rename the columns to match the R solution
result_tibble.columns = ['average_gain', 'name']

# Display the resulting tibble
print(result_tibble)

```

	average_gain	name
0	1.285235	AirTran Airways Corporation
1	0.545455	Frontier Airlines Inc.
2	0.063590	Envoy Air
3	-0.909774	Hawaiian Airlines Inc.
4	-2.837182	US Airways Inc.
5	-2.846154	SkyWest Airlines Inc.
6	-3.017921	Mesa Airlines Inc.
7	-4.905070	JetBlue Airways
8	-5.787007	ExpressJet Airlines Inc.
9	-7.960461	Delta Air Lines Inc.
10	-8.571100	Southwest Airlines Co.
11	-8.669899	American Airlines Inc.
12	-9.621162	United Air Lines Inc.
13	-9.875375	Virgin America
14	-10.288170	Endeavor Air Inc.
15	-22.211409	Alaska Airlines Inc.

```

[44]: # Read the transactions.csv dataset
transactions = pd.read_csv("transactions.csv")

# Reshape the data to be tidy
transactions = pd.melt(transactions, id_vars=['person', 'month'],
↳var_name='transaction_type', value_name='amount')

# Filter the data for Jenna or John
transactions = transactions[transactions['person'].isin(['jenna', 'john'])]

```

```

# Get the unique months in which either Jenna or John had a transaction
unique_months = transactions['month'].unique()

# Create a pandas DataFrame with the unique months
result_tibble = pd.DataFrame({'month': unique_months})

# Display the resulting tibble
print(result_tibble)

```

```

      month
0         1
1         3

```

1.5 Question 5 (25 points)

Solve again exercises 1 and 2 of Homework 3, but this time use `pandas`!

You are encouraged to use the published solutions of Homework 3 and “translate” them from R’s `tidyverse`/`dplyr` to Python’s `pandas`.

Again, you might find this [“translation guide”](#) helpful!

```

[53]: # Read the datasets
dinners = pd.read_csv("dinners.csv")
drinks = pd.read_csv("drinks.csv")
food = pd.read_csv("food.csv")
persons = pd.read_csv("persons.csv")

# Merge the datasets to create the dinners_explicit tibble
dinners_explicit = pd.merge(dinners, drinks, left_on="drink_id",
    ↪right_on="item_id", how="left")
dinners_explicit = dinners_explicit.rename(columns={"price": "drink_price",
    ↪"item_name": "drink"})

dinners_explicit = pd.merge(dinners_explicit, food, left_on="food_id",
    ↪right_on="food_id", how="left")
dinners_explicit = dinners_explicit.rename(columns={"price": "food_price",
    ↪"name": "food"})

dinners_explicit = pd.merge(dinners_explicit, persons, left_on="person_id",
    ↪right_on="id", how="left")

# Select the specified columns
dinners_explicit = dinners_explicit[["drink", "drink_price", "food",
    ↪"food_price", "first_name", "last_name", "age"]]

# Display the resulting tibble

```

```
dinners_explicit
```

```
[53]:      drink drink_price      food food_price first_name last_name \
0      NaN      NaN      pasta    $8.50      Valter  Evangelista
1      NaN      NaN  ice cream    $4.50      Polly      Verity
2      water    $1.00      NaN      NaN      NaN      NaN
3      beer    $5.00      cake    $4.50      Aysha      Freitas
4      NaN      NaN      pizza    $12      Rayno      Van Kann
5      water    $1.00      fish    $15.00     Valter  Evangelista
6      NaN      NaN      pizza    $12      Rayno      Van Kann
7  sparkling water    $2.00  ice cream    $4.50     Ksenya      Dunai
8      soda    $2.50  pop corn    $1.50      Polly      Verity
9      water    $1.00      salad    $5.00      Aysha      Freitas
10     wine    $9.00      NaN      NaN      NaN      NaN
11     soda    $2.50      salad    $5.00      NaN      NaN
12     beer    $5.00      cake    $4.50      Aysha      Freitas
13     NaN      NaN      steak    $12.00     NaN      NaN
14     wine    $9.00      NaN      NaN     Valter  Evangelista
15     wine    $9.00  pop corn    $1.50      NaN      NaN
16     NaN      NaN      fish    $15.00     NaN      NaN
17     water    $1.00      fries    $3.00      Polly      Verity
18  sparkling water    $2.00  burger    $5.00      Polly      Verity
19     soda    $2.50      steak    $12.00     Polly      Verity
20     NaN      NaN      fries    $3.00     Valter  Evangelista
21     NaN      NaN      NaN      NaN      NaN      NaN
22     NaN      NaN      pasta    $8.50      NaN      NaN
23     soda    $2.50      pasta    $8.50      NaN      NaN
24  sparkling water    $2.00      salad    $5.00      NaN      NaN
25  sparkling water    $2.00      pizza    $12      Rayno      Van Kann
26  sparkling water    $2.00  ice cream    $4.50      NaN      NaN
27  sparkling water    $2.00      fish    $15.00     NaN      NaN
28     NaN      NaN      fries    $3.00      Polly      Verity
29     beer    $5.00      pasta    $8.50     Valter  Evangelista
30     water    $1.00  burger    $5.00      NaN      NaN
31     NaN      NaN      pasta    $8.50      Rayno      Van Kann
32     NaN      NaN      salad    $5.00      Aysha      Freitas
33     NaN      NaN      fries    $3.00      NaN      NaN
34     soda    $2.50      cake    $4.50      NaN      NaN
35  sparkling water    $2.00      fish    $15.00     NaN      NaN
36     soda    $2.50      NaN      NaN      NaN      NaN
37     beer    $5.00  pop corn    $1.50      NaN      NaN
38     wine    $9.00      fish    $15.00     NaN      NaN
39     NaN      NaN      salad    $5.00      NaN      NaN
40     soda    $2.50      NaN      NaN      NaN      NaN
41     NaN      NaN      fries    $3.00      NaN      NaN
42     wine    $9.00      pizza    $12      Polly      Verity
43     soda    $2.50      NaN      NaN      Aysha      Freitas
```


44	water	\$1.00	pizza	\$12	NaN	NaN
45	water	\$1.00	fries	\$3.00	Polly	Verity
46	NaN	NaN	pop corn	\$1.50	NaN	NaN
47	soda	\$2.50	pizza	\$12	Rayno	Van Kann
48	sparkling water	\$2.00	pizza	\$12	Aysha	Freitas
49	NaN	NaN	steak	\$12.00	Polly	Verity

	age
0	34.0
1	61.0
2	NaN
3	55.0
4	29.0
5	34.0
6	29.0
7	31.0
8	61.0
9	55.0
10	NaN
11	NaN
12	55.0
13	NaN
14	34.0
15	NaN
16	NaN
17	61.0
18	61.0
19	61.0
20	34.0
21	NaN
22	NaN
23	NaN
24	NaN
25	29.0
26	NaN
27	NaN
28	61.0
29	34.0
30	NaN
31	29.0
32	55.0
33	NaN
34	NaN
35	NaN
36	NaN
37	NaN
38	NaN

```
39 NaN
40 NaN
41 NaN
42 61.0
43 55.0
44 NaN
45 61.0
46 NaN
47 29.0
48 55.0
49 61.0
```

```
[81]: # Assuming you have the 'dinners_explicit' DataFrame from the previous steps
```

```
# Create a summary tibble for the two most popular drinks
top2_drinks = dinners_explicit['drink'].value_counts().reset_index()
top2_drinks.columns = ['drink', 'n']
top2_drinks = top2_drinks.head(2)

# Create a summary tibble for the two most popular foods
top2_foods = dinners_explicit['food'].value_counts().reset_index()
top2_foods.columns = ['food', 'n']
top2_foods = top2_foods.head(2)

# Display the summary tibbles
print("Top 2 Drinks:")
top2_drinks
```

Top 2 Drinks:

```
[81]:      drink  n
0      soda   9
1  sparkling water  8
```

```
[82]: print("\nTop 2 Foods:")
      top2_foods
```

Top 2 Foods:

```
[82]:    food  n
0  pizza  7
1  fries  6
```

```
[5]: #!pip install nbconvert
```